

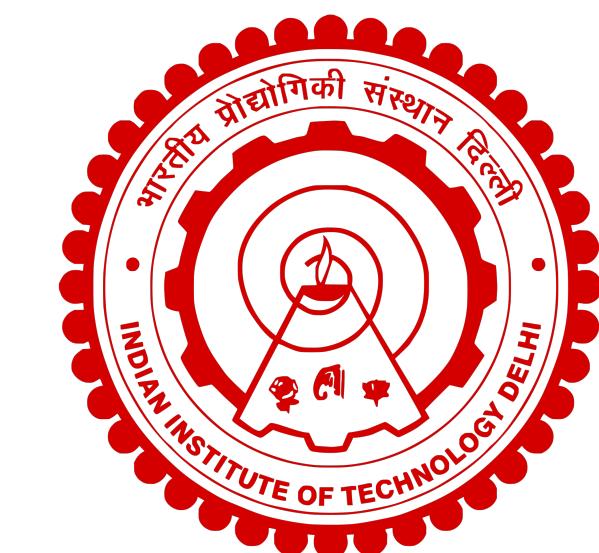
COL 362/632

Introduction To Database Management Systems

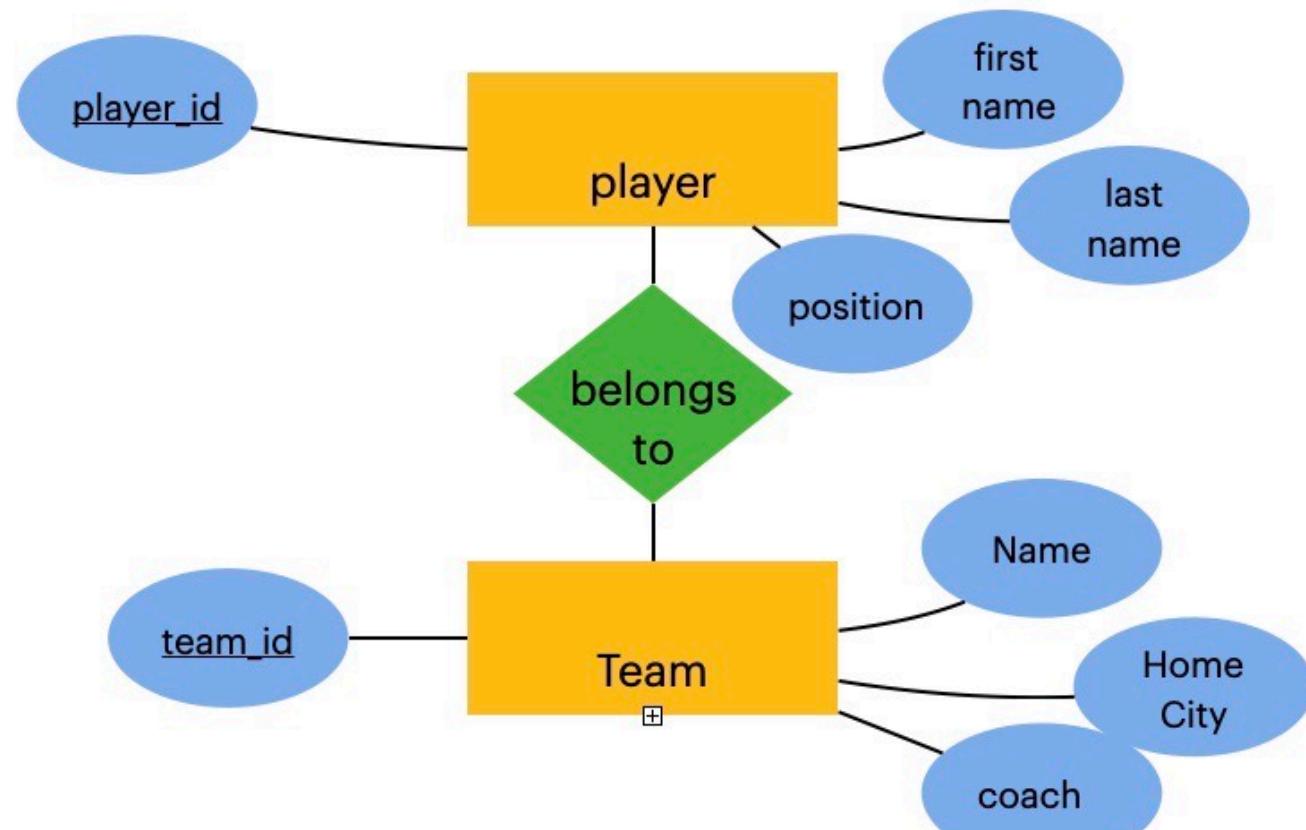
Database Internals – Overview

Kaustubh Beedkar

Computer Science & Engineering
Indian Institute of Technology Delhi



Why Databases are cool



```
%%sql SELECT category.name, AVG(film.rental_duration)
AS average_duration_category,
(SELECT AVG(rental_duration)
FROM film
) AS overall_average_duration
FROM category
JOIN film_category ON category.id = film_category.category_id
JOIN film ON film_category.film_id = film.id
GROUP BY category.name
```

Really?



Normal Forms

Player

PlayerId	Name	Position	
123	MS Dhoni	7	
193	Virat Kohli	3	
131	Rohit Sharma	1	
134	SK Yadav	5	MI
154	David Warner	3	DC

Boyce-Codd normal form (BCNF)

A relation R is in BCNF iff for every non-trivial FD $X \rightarrow Y$, the conditions hold

- ▶ $X \rightarrow Y$ is trivial FD
- ▶ X is a superkey

Normal Form

Third Normal Form (3NF)

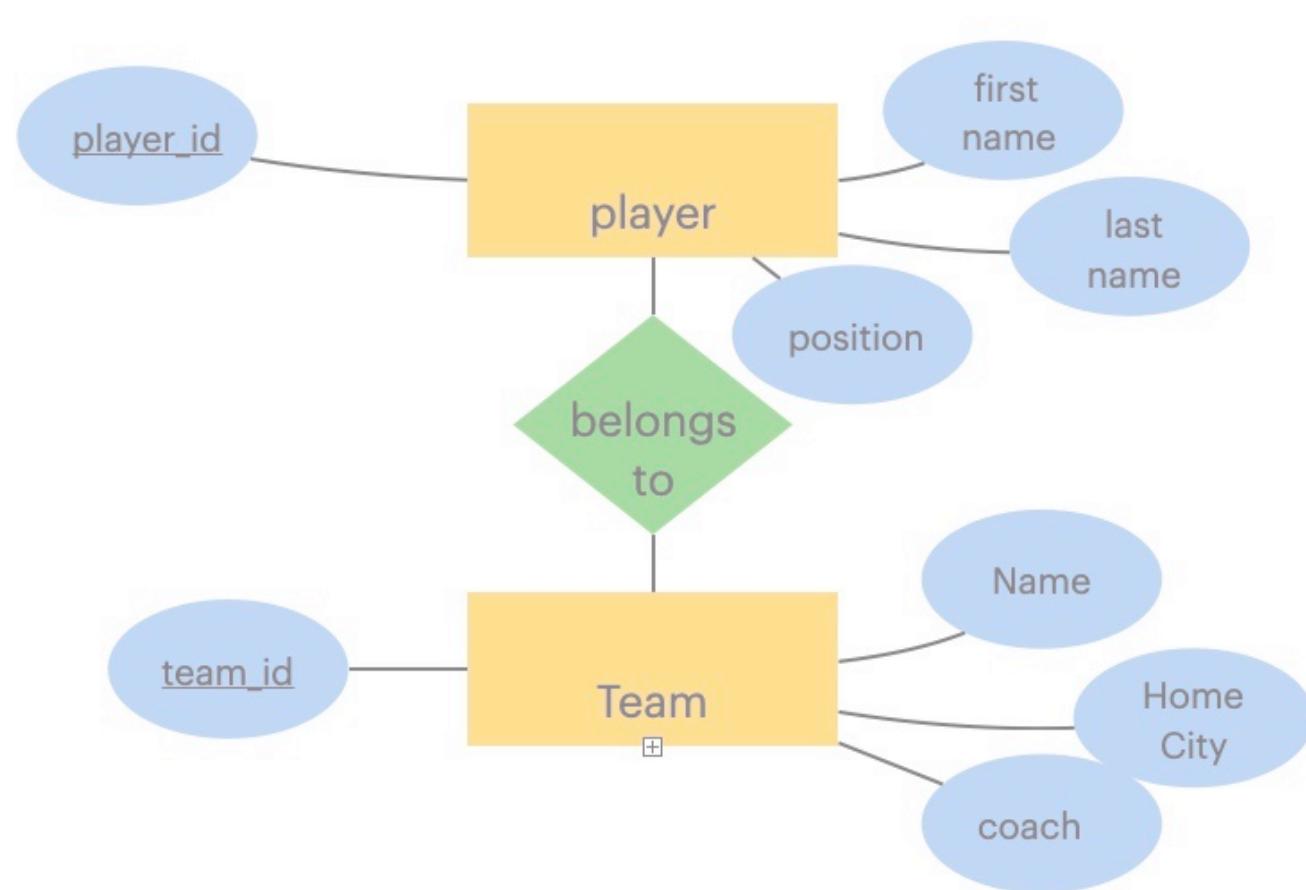
A relation R is in 3NF if whenever $X \rightarrow Y$, one of the following is true

1. $X \rightarrow Y$ is trivial FD
2. X is superkey
3. $\forall a \in Y \setminus X, a$ is prime attribute



Note: If R is in BCNF \implies it is in 3NF

Why Databases are cool



Normal Forms

Player			Boyce-Codd normal form (BCNF)
PlayerId	Name	Position	PlaysFor
123	MS Dhoni	7	CSK
193	Virat Kohli	3	RCB
131	Rohit Sharma	1	MI
134	SK Yadav	5	MI
154	David Warner	3	DC

A relation R is in BCNF iff for every FD $X \rightarrow Y$, at least one of:
 $\triangleright X \rightarrow Y$ is trivial FD
 $\triangleright X$ is a superkey

Normal Forms

Third Normal Form (3NF)

A relation R is in 3NF if whenever $X \rightarrow Y$, one of the following is true

1. $X \rightarrow Y$ is trivial FD
2. X is superkey
3. $\forall a \in Y \setminus X, a$ is prime attribute

Note: If R is in BCNF \implies it is in 3NF

```

%%sql SELECT category.name, AVG(film.rental_duration)
AS average_duration_category,
(SELECT AVG(rental_duration)
FROM film
) AS overall_average_duration
FROM category
JOIN film_category
JOIN film ON film_category.film_id = film.film_id
GROUP BY category.name
  
```

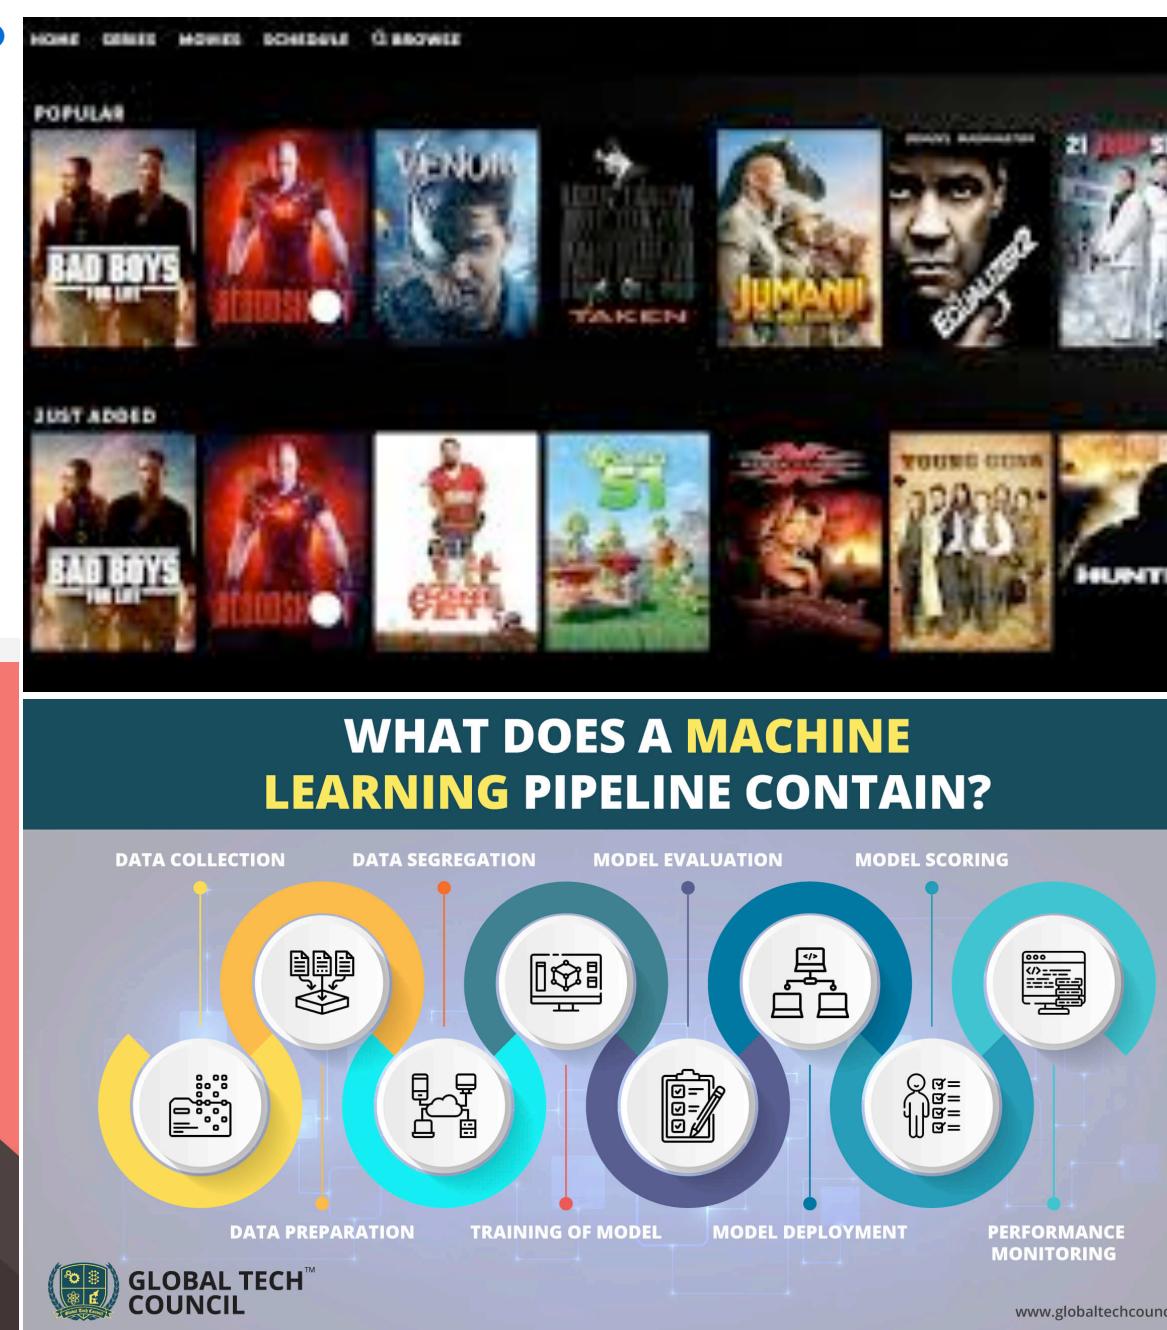
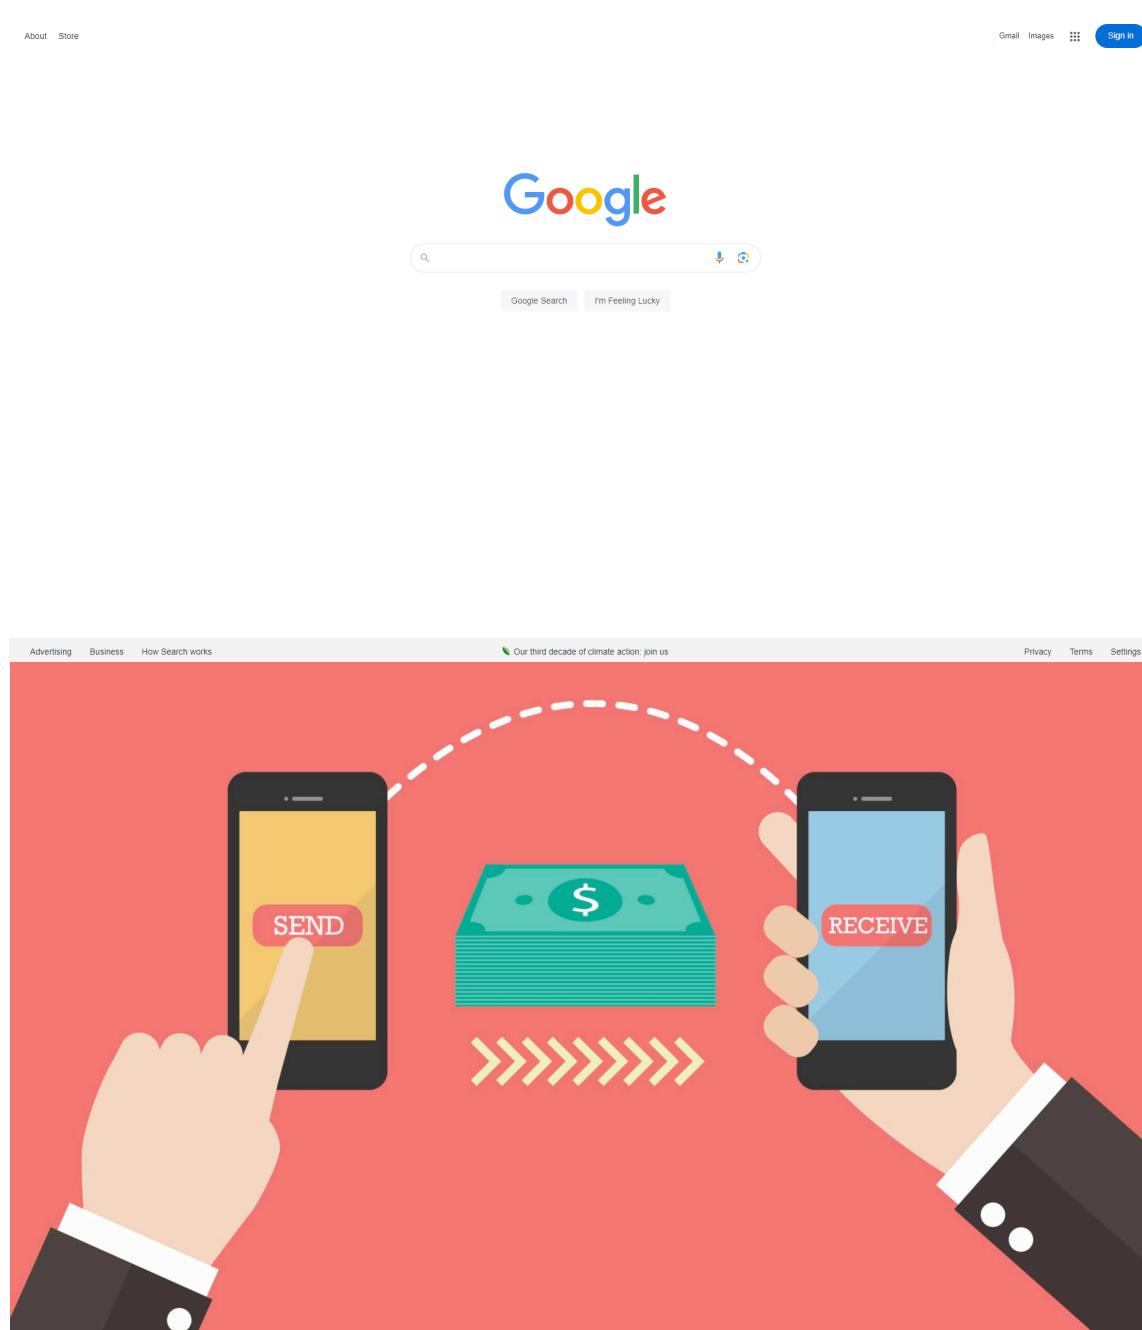
This is ~~cool~~ **important** for

- User of the database
- Application developer

Not so exciting as a computer scientist!

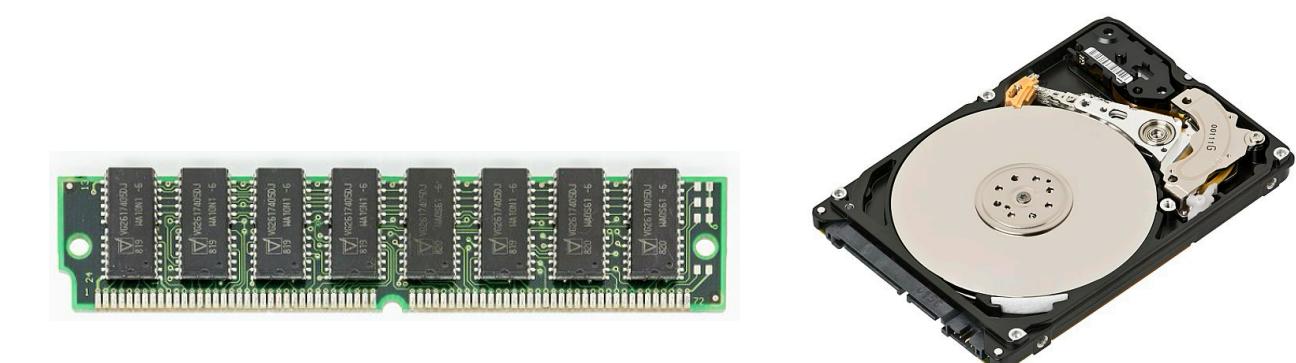
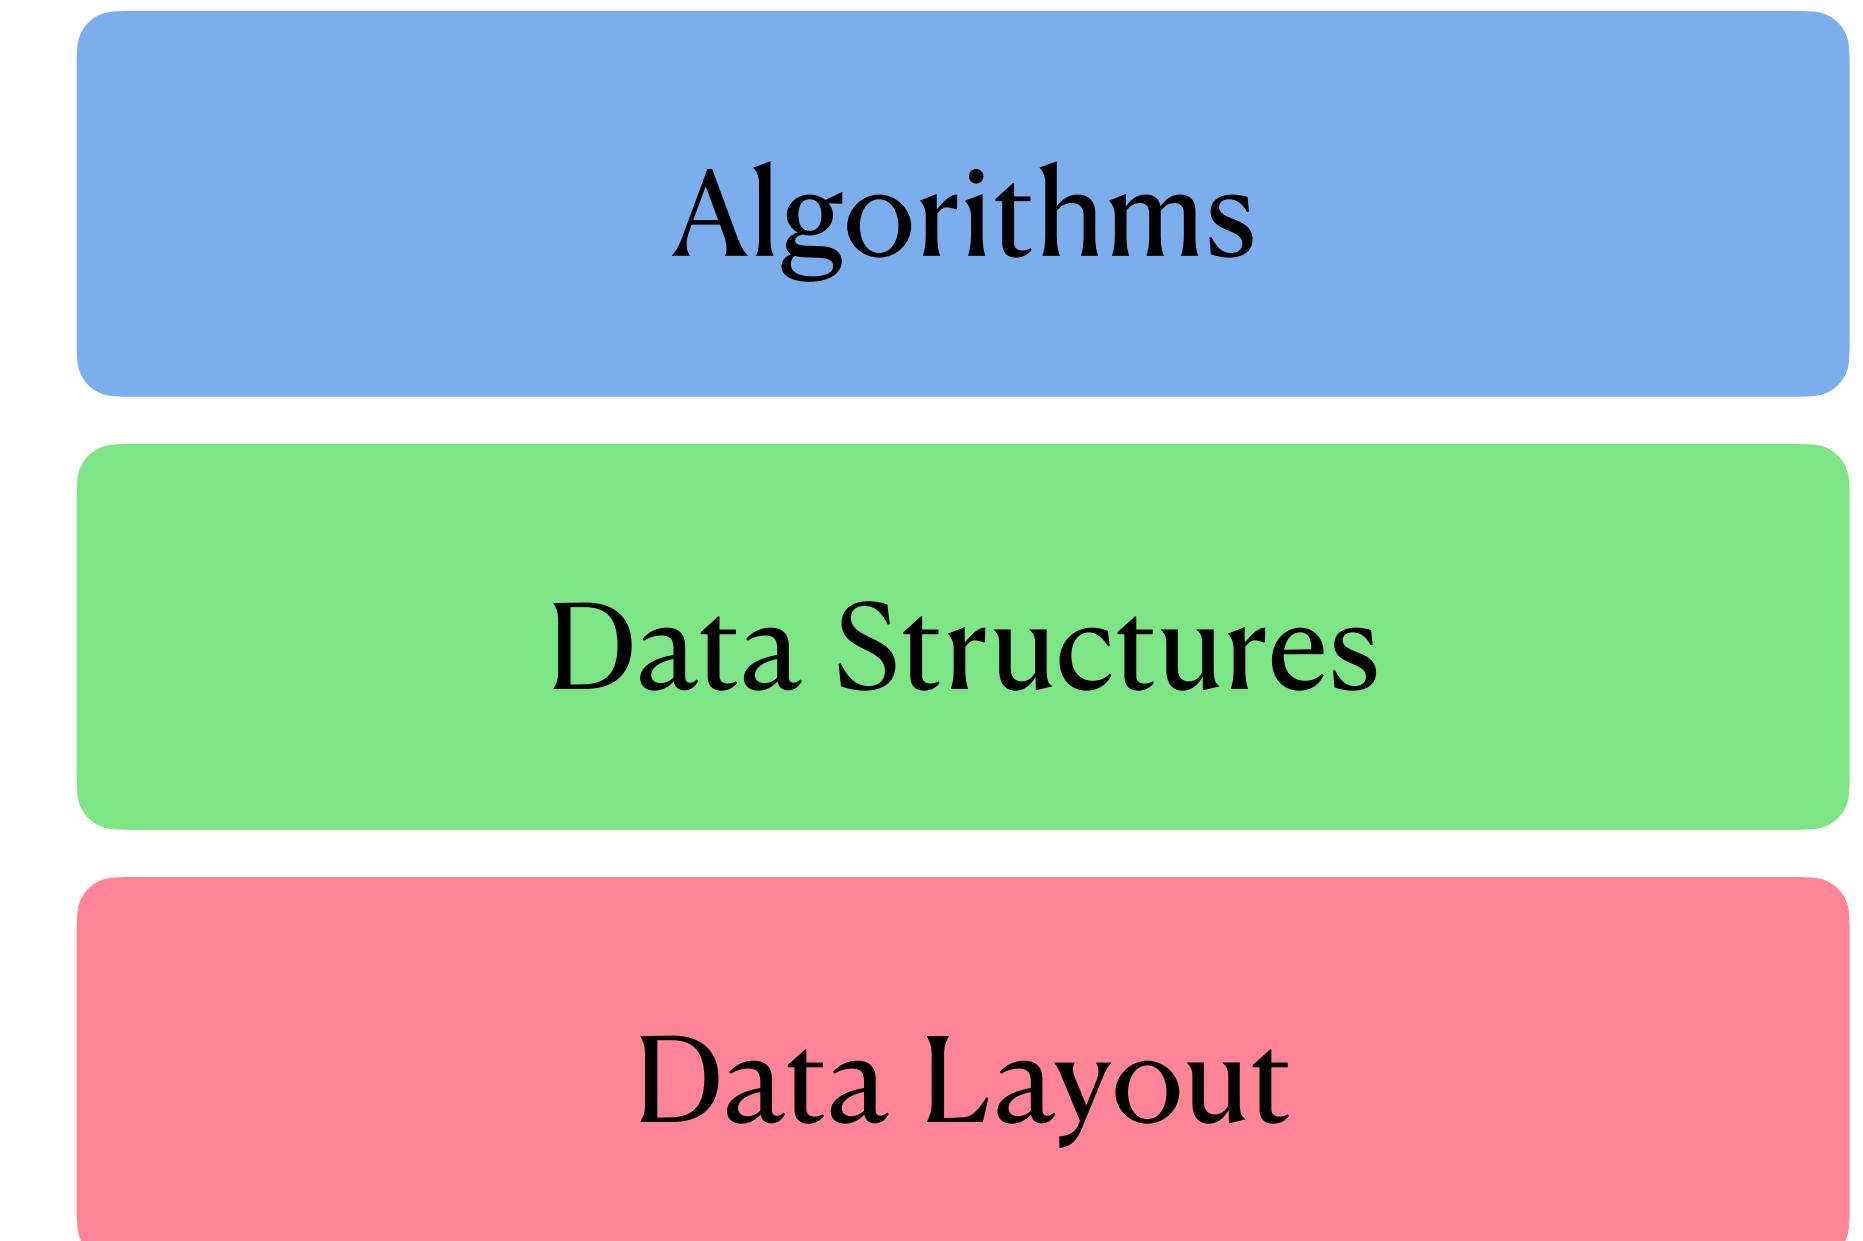
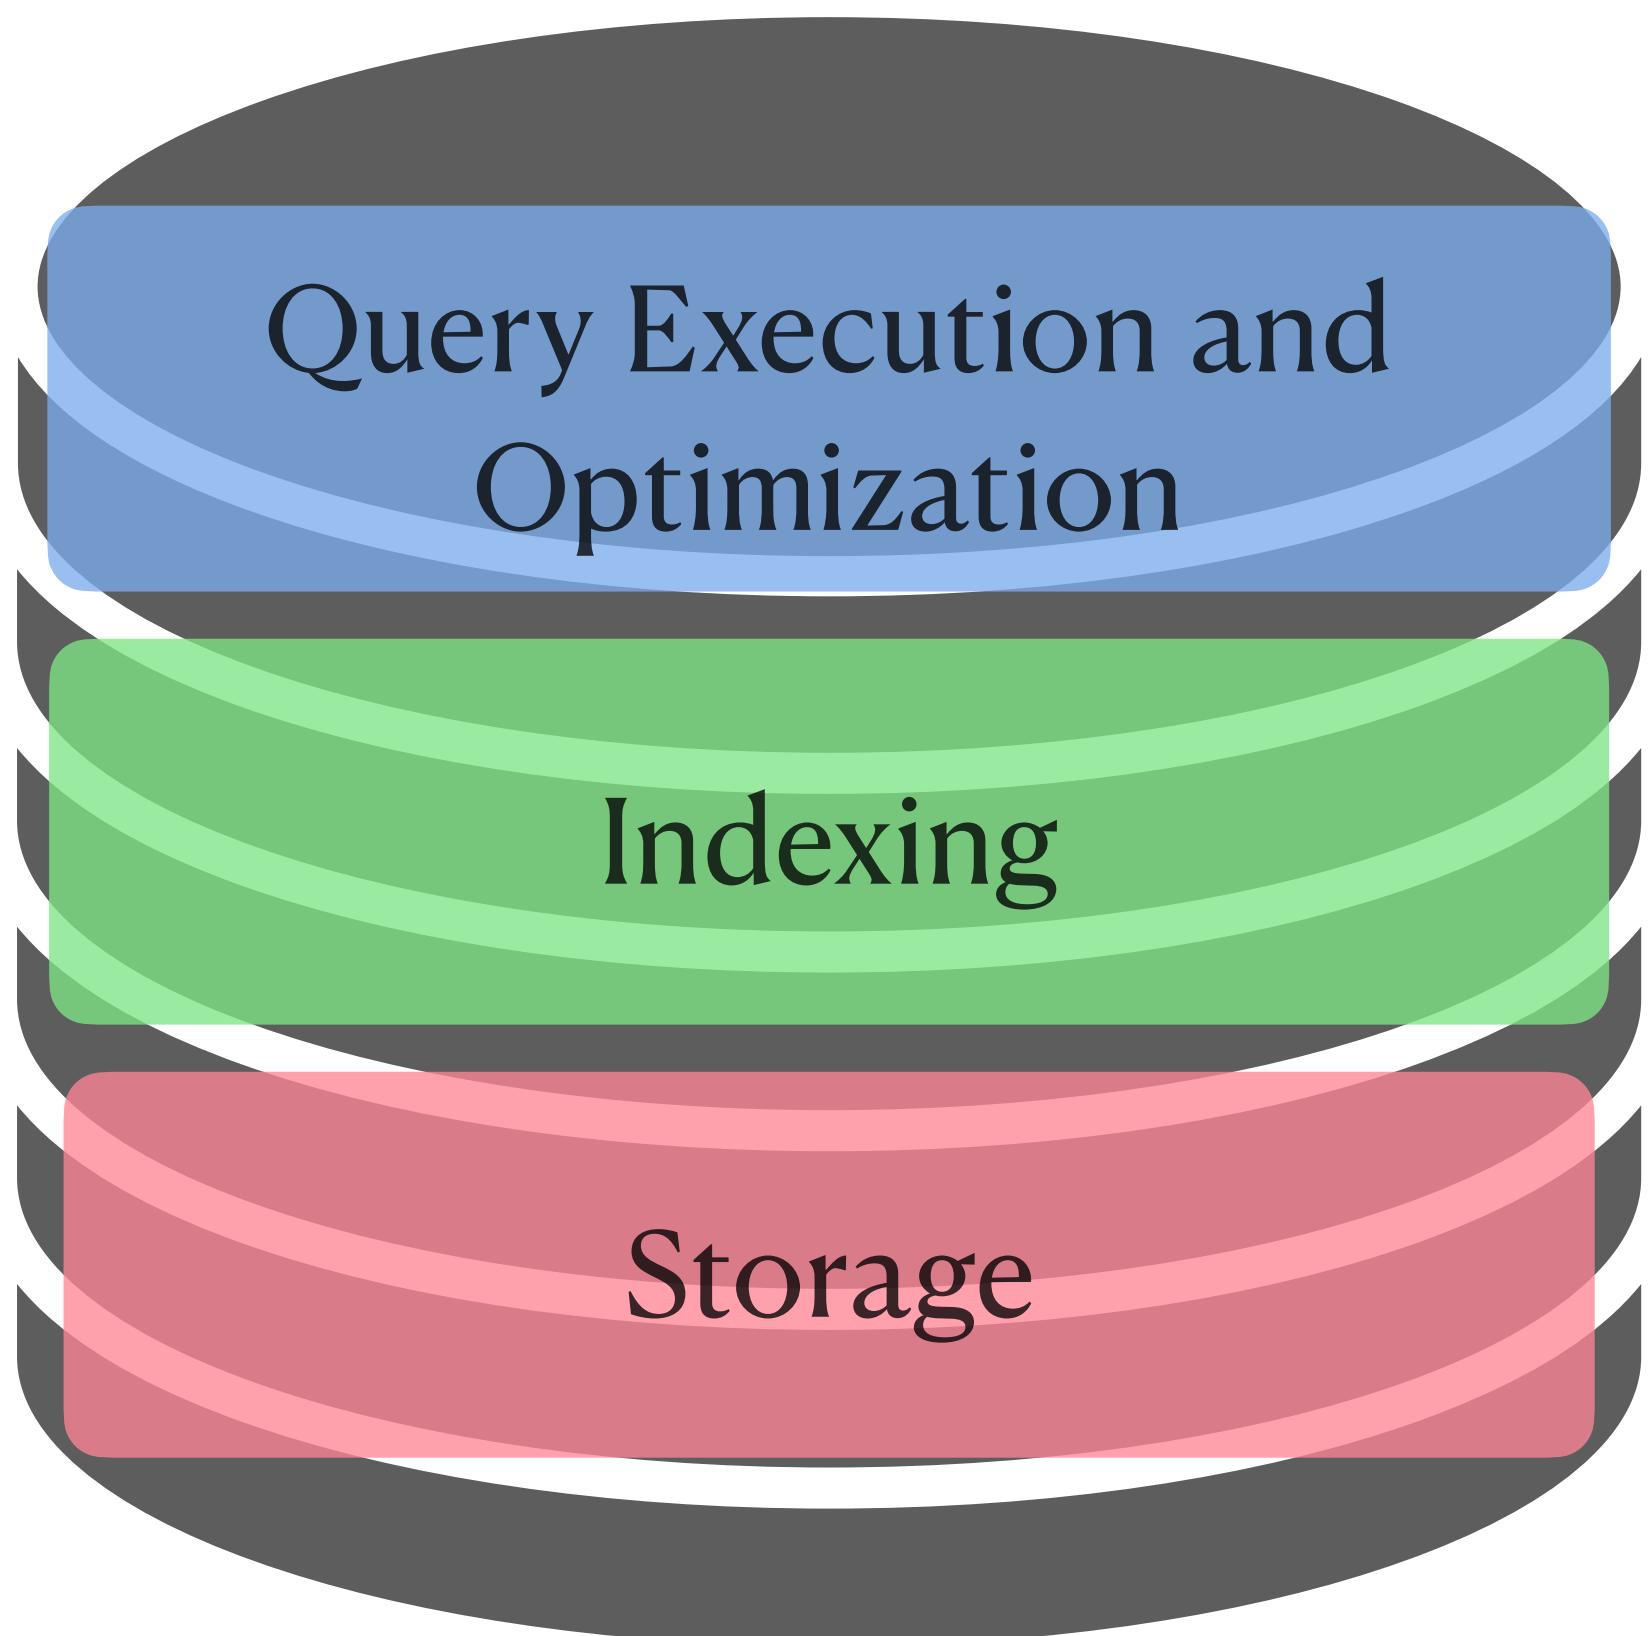
Why Databases are cool

- Database **technology** is cool!
- Powering the digital world

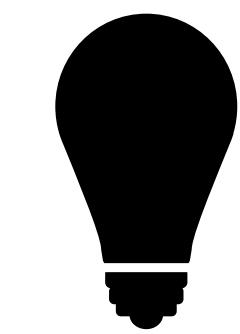


DB internals

- Three fundamental layers



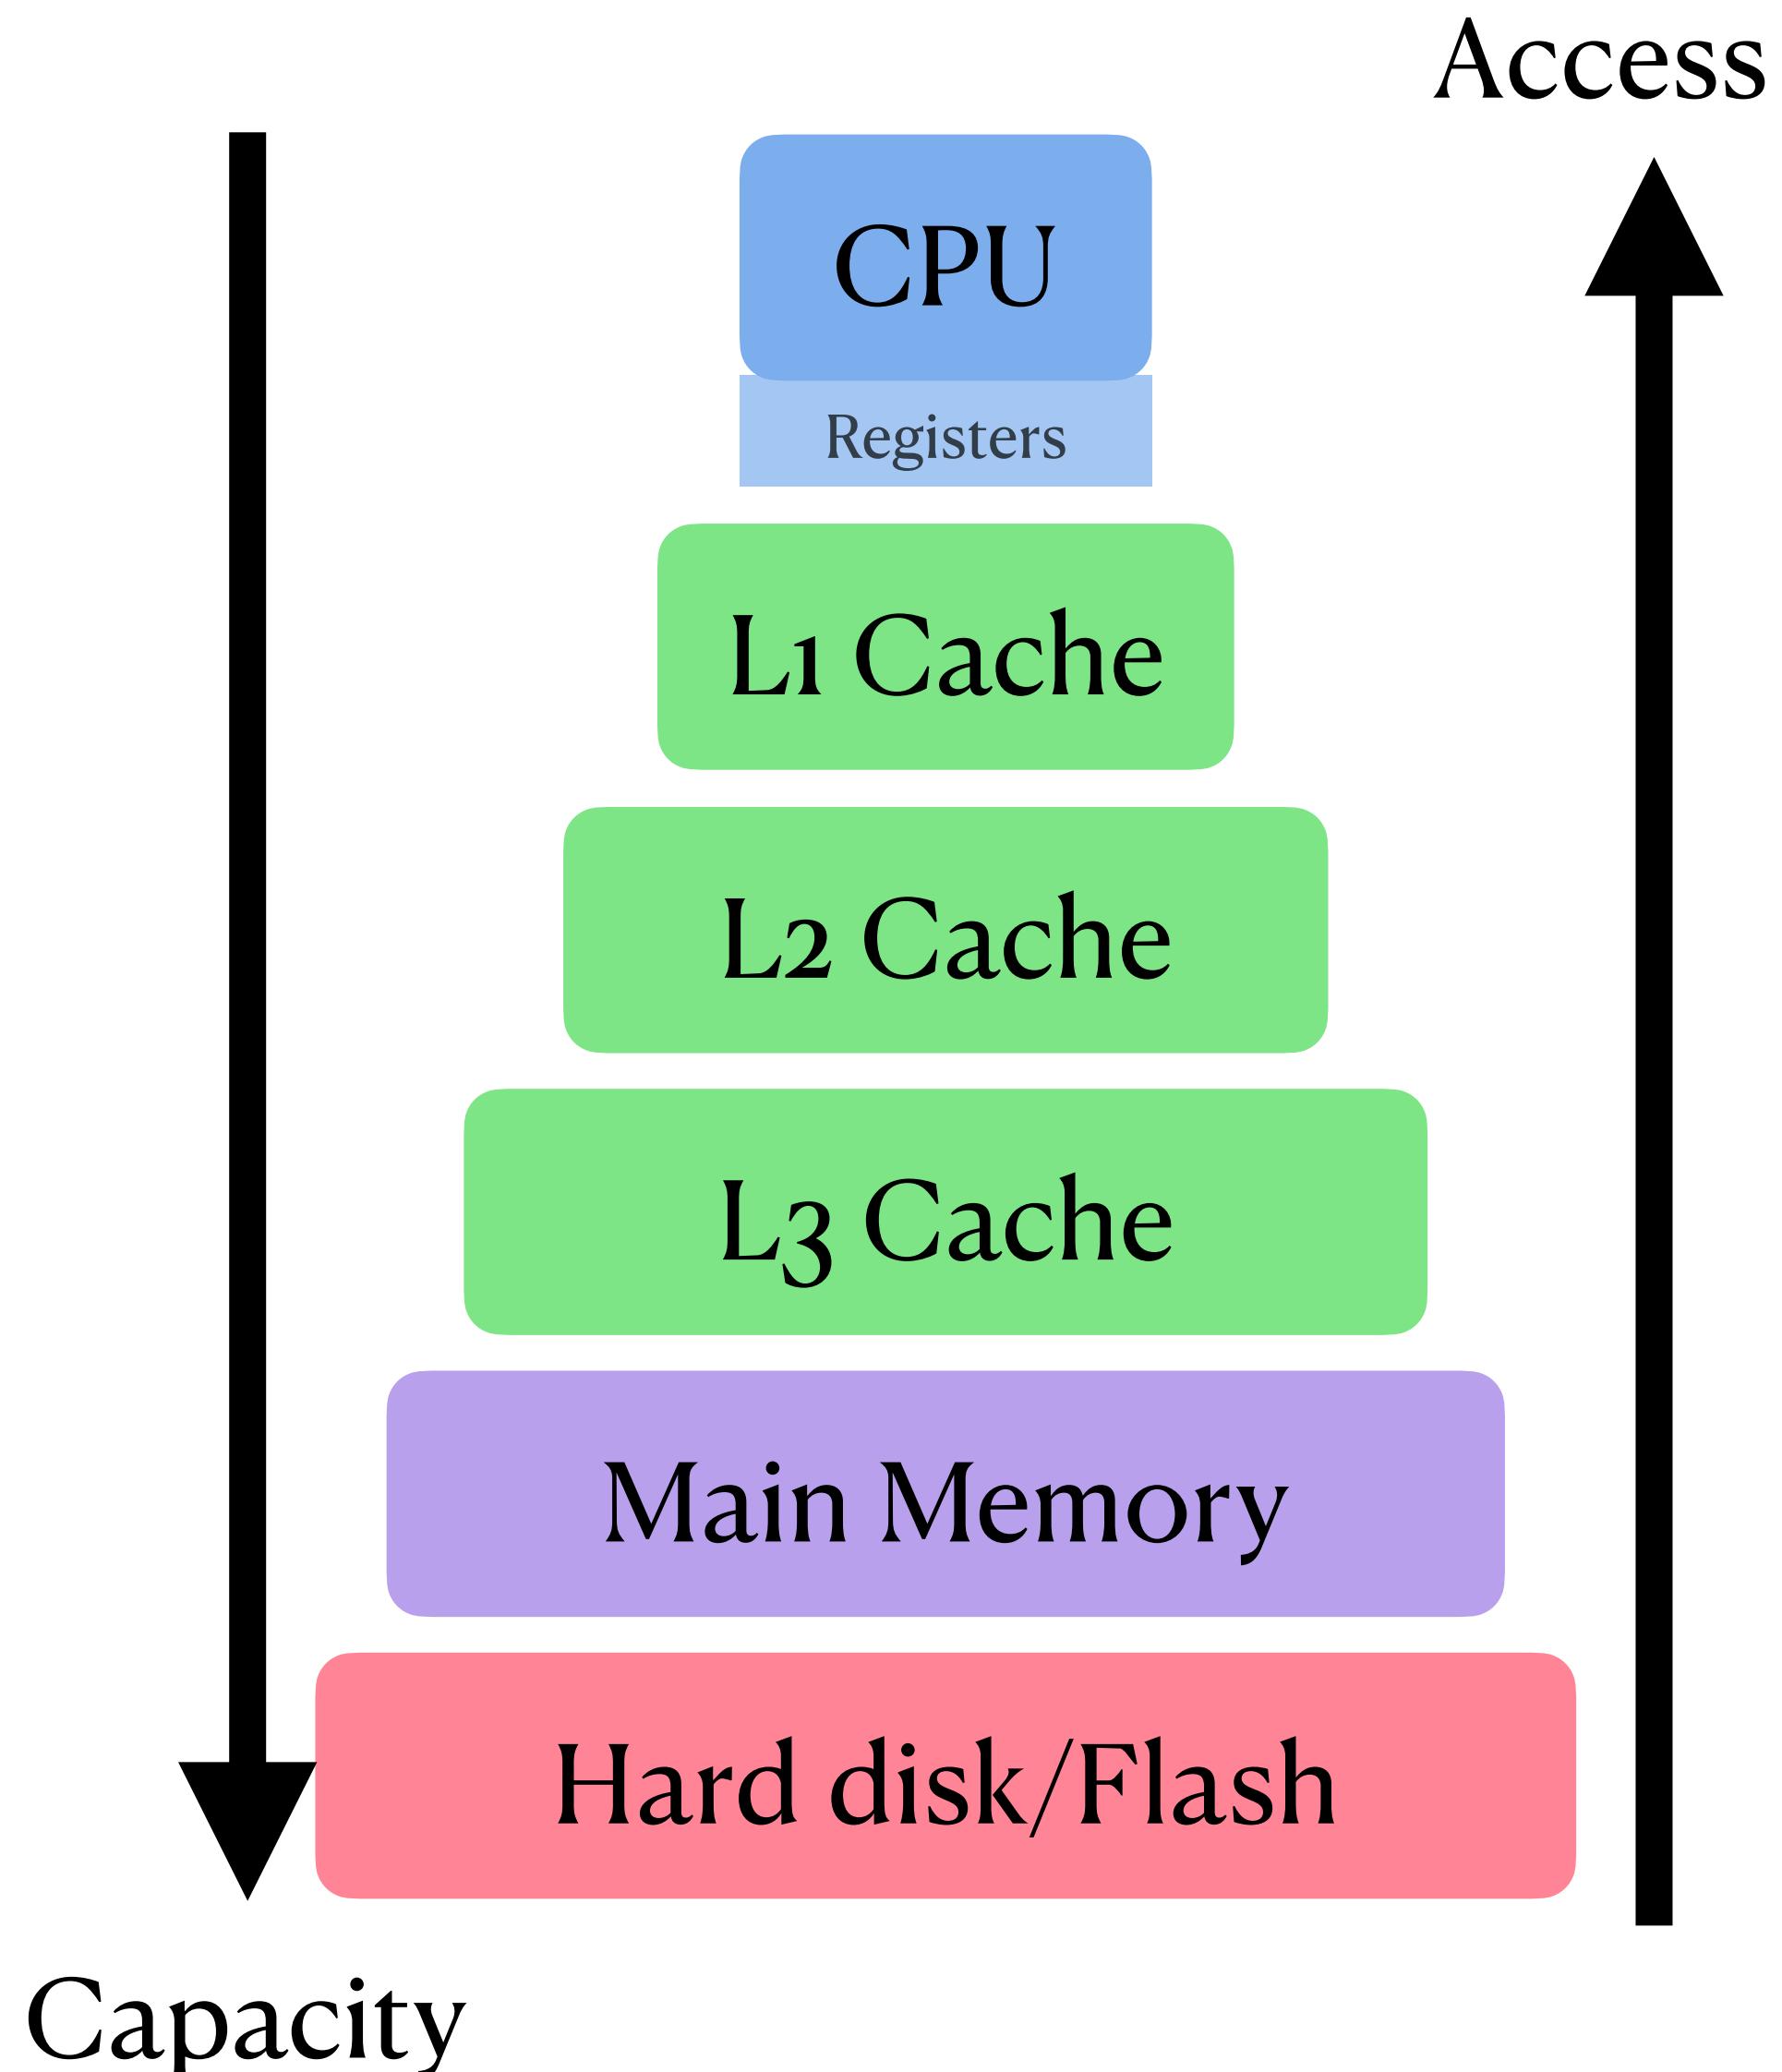
Storage



Imagine you have to store a trillion rows of data and retrieve them efficiently. How would you do it?

Challenges of Storage

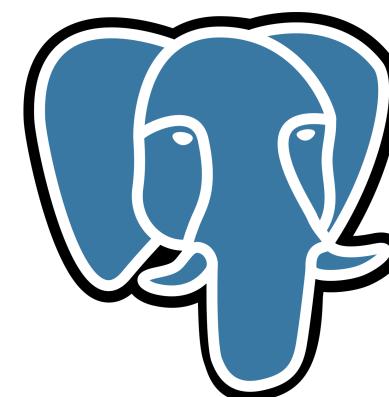
- Storing huge amounts of data while keeping it **durable, fault-tolerant, and efficient**
- Dealing with disk vs. memory trade-offs.



Storage

Key Concepts

- Buffer management
- Storage structures
 - Heap files: unstructured simple storage
 - B-Trees: used in most relational databases
 - LSM Trees: log-structured merged trees, used in No-SQL databases
 - Column Stores: used in analytical databases Snowflake, Google BigQuery



Storage

Key Concepts

- Data layout

Player			
PlayerId	Name	Position	PlaysFor
123	MS Dhoni	7	CSK
193	Virat Kohli	3	RCB
131	Rohit Sharma	1	MI
134	SK Yadav	5	MI
154	David Warner	3	DC



2-Dimentional

1-Dimentional

Player

PlayerId	Name	Position	PlaysFor
123	MS Dhoni	7	CSK
193	Virat Kohli	3	RCB
131	Rohit Sharma	1	MI
134	SK Yadav	5	MI
154	David Warner	3	DC

Row layout

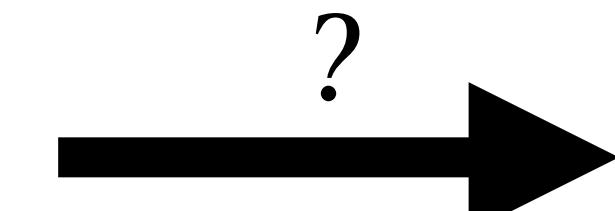
...123 | MS Dhoni | 7 | CSK || 193 | Virat Kohli | 3 | RCB || 131 | Rohit Sharma | 1 | MI || ...

Storage

Key Concepts

- Data layout

Player			
PlayerId	Name	Position	PlaysFor
123	MS Dhoni	7	CSK
193	Virat Kohli	3	RCB
131	Rohit Sharma	1	MI
134	SK Yadav	5	MI
154	David Warner	3	DC



2-Dimentional

1-Dimentional

Player

PlayerId	Name	Position	PlaysFor
123	MS Dhoni	7	CSK
193	Virat Kohli	3	RCB
131	Rohit Sharma	1	MI
134	SK Yadav	5	MI
154	David Warner	3	DC

Column layout

123 | 193 | 131 | 134 | 154

|| MS Dhoni | Virat Kolhi | Rohit Sharma | SK Yadav | David..

Storage

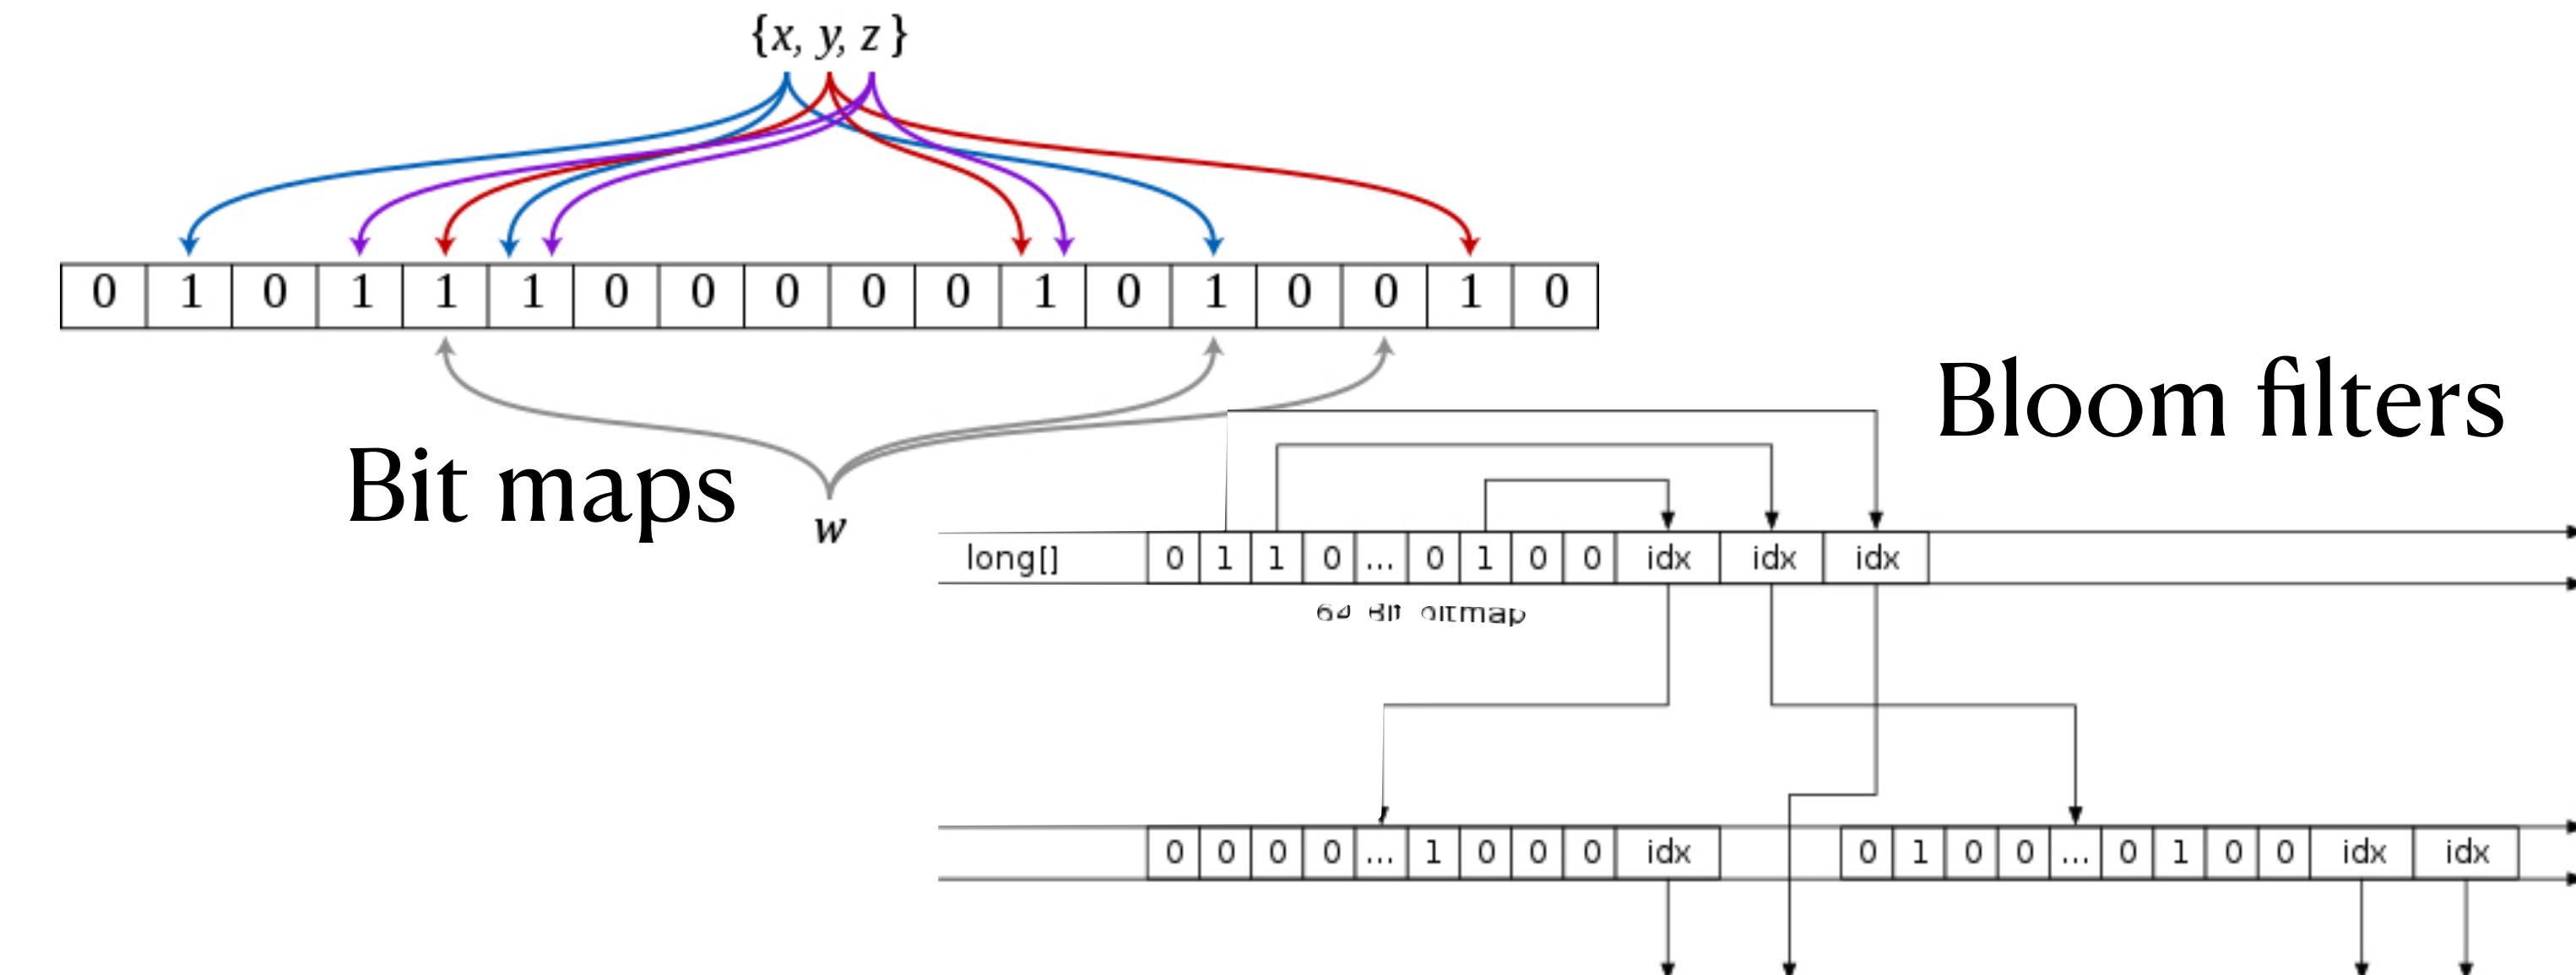
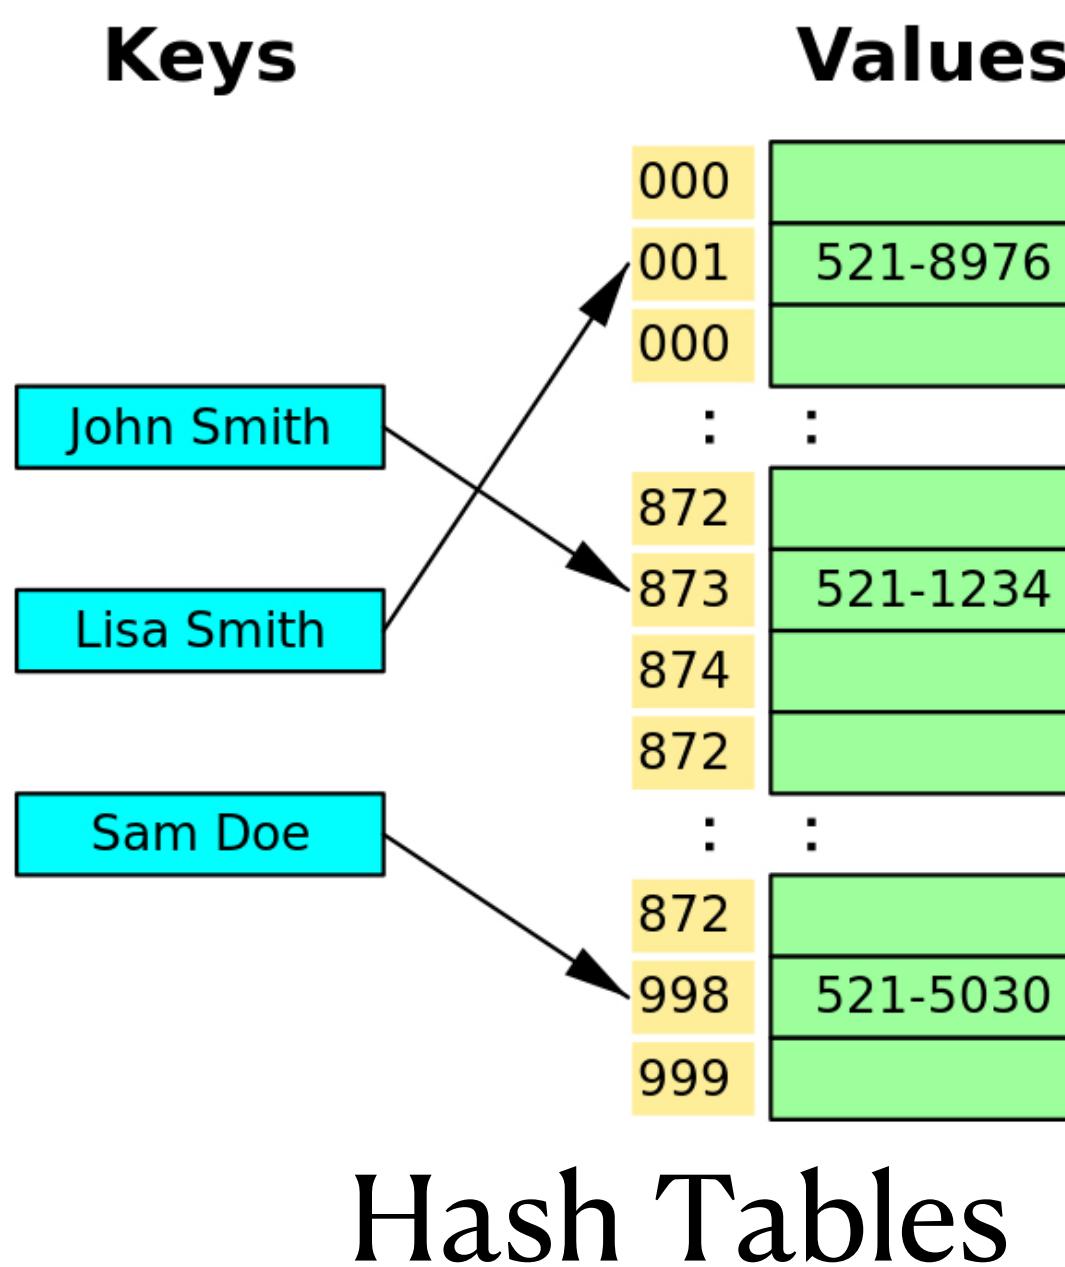
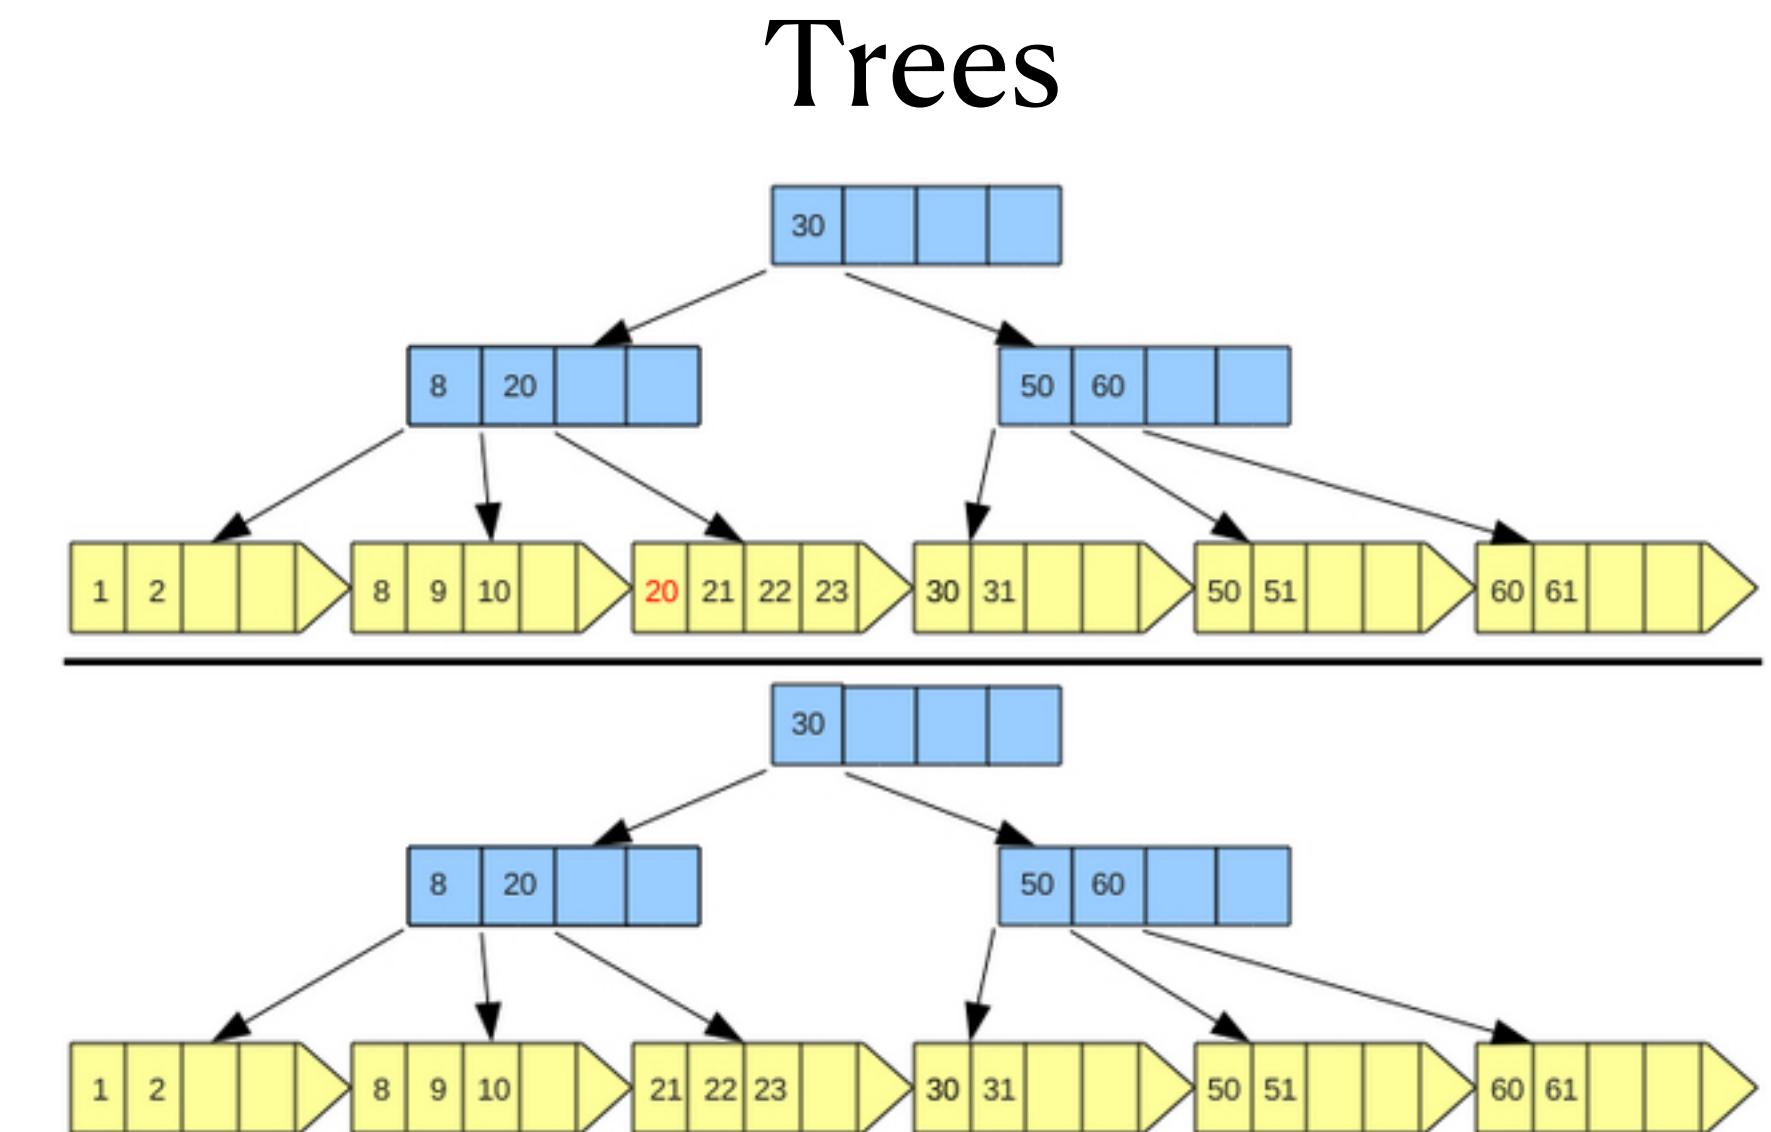
Beyond Databases

- Cloud storage
- Blockchain
- High performance logging



Indexing

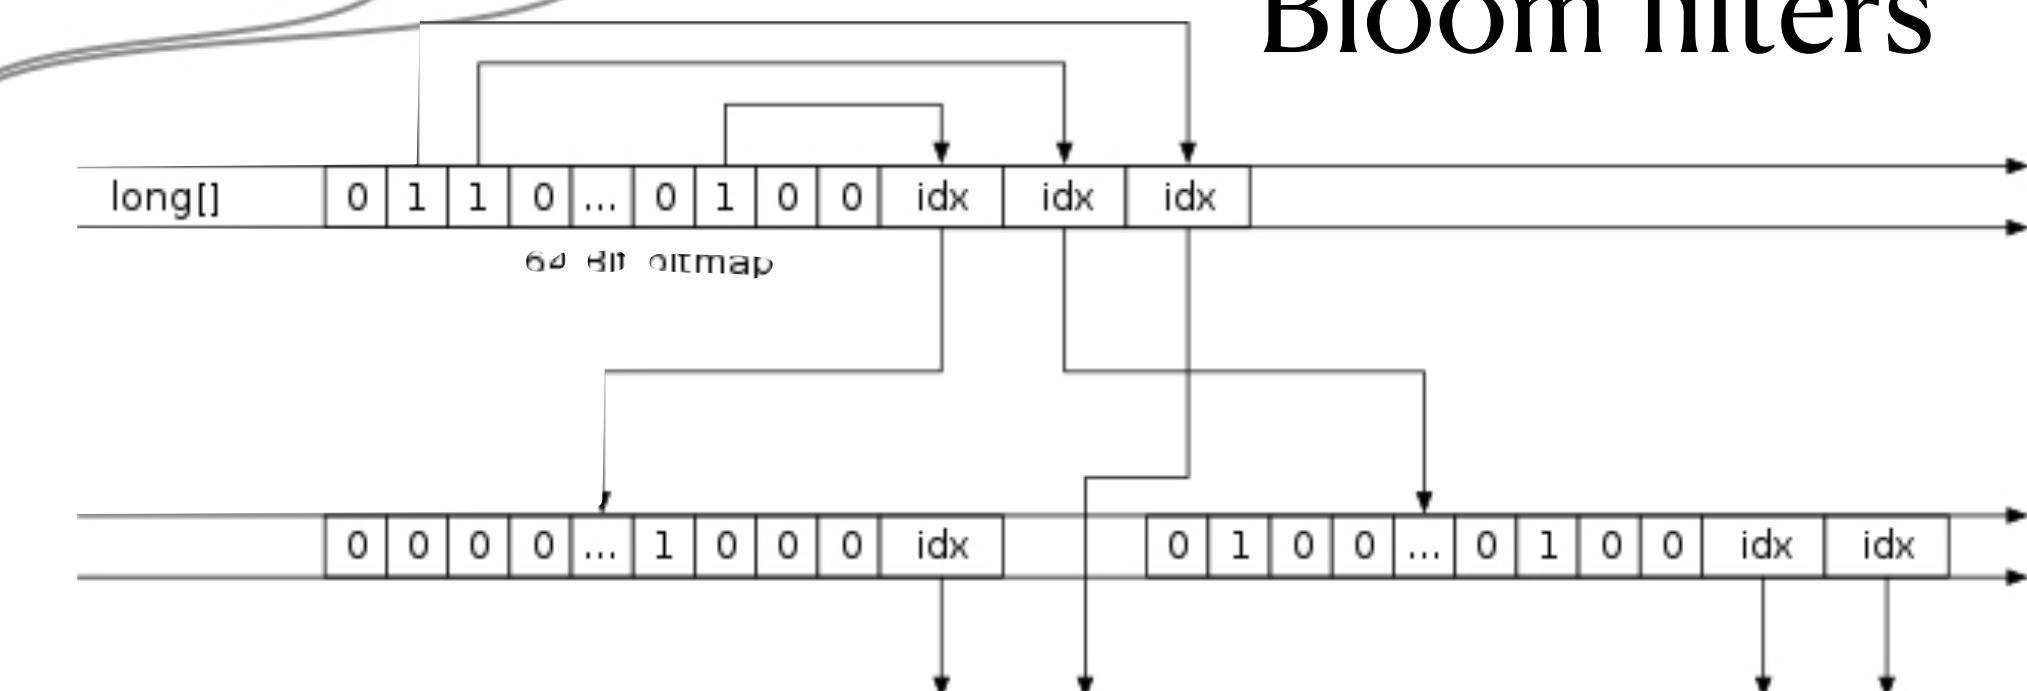
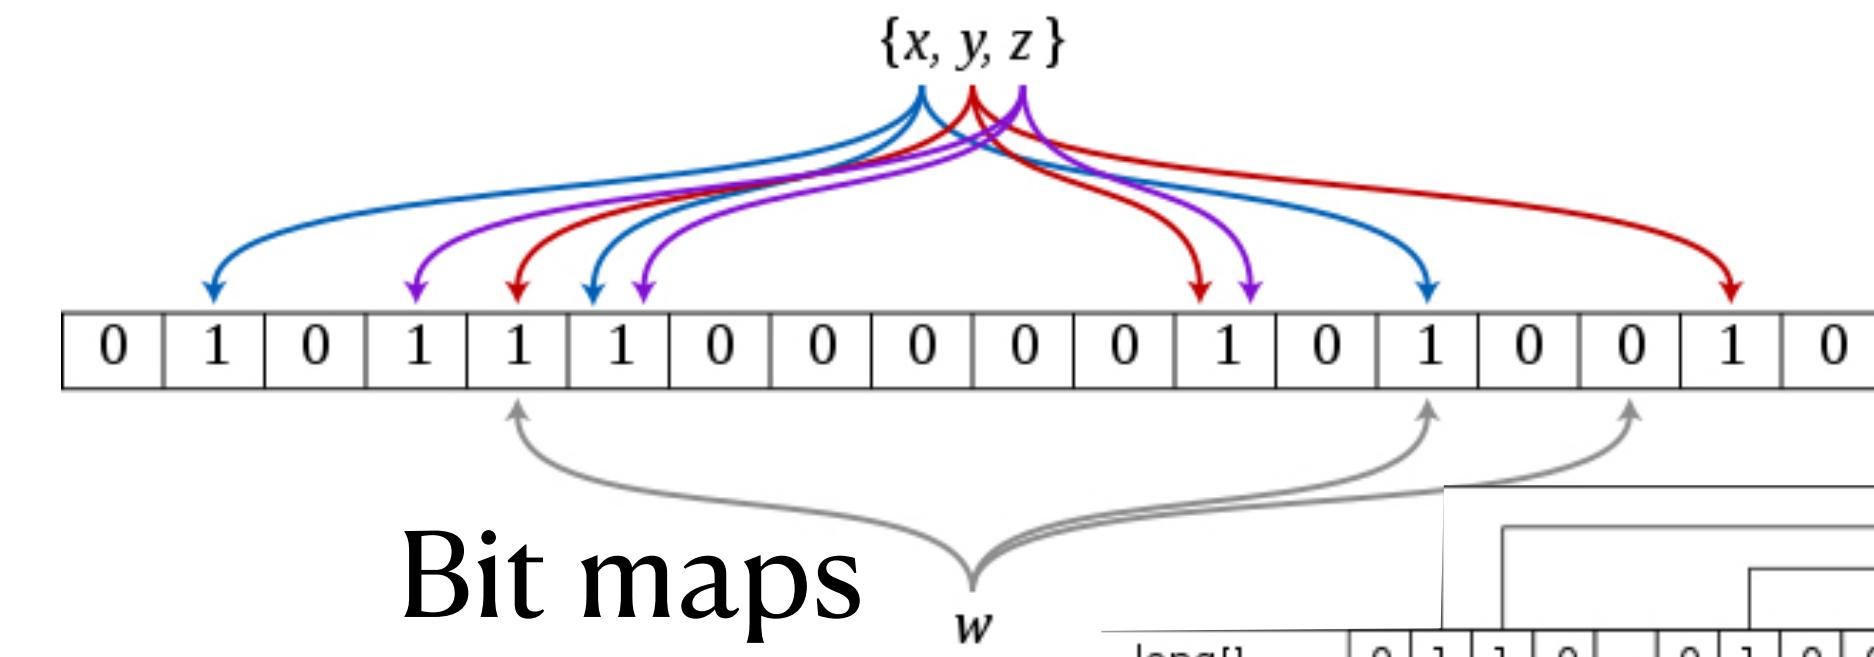
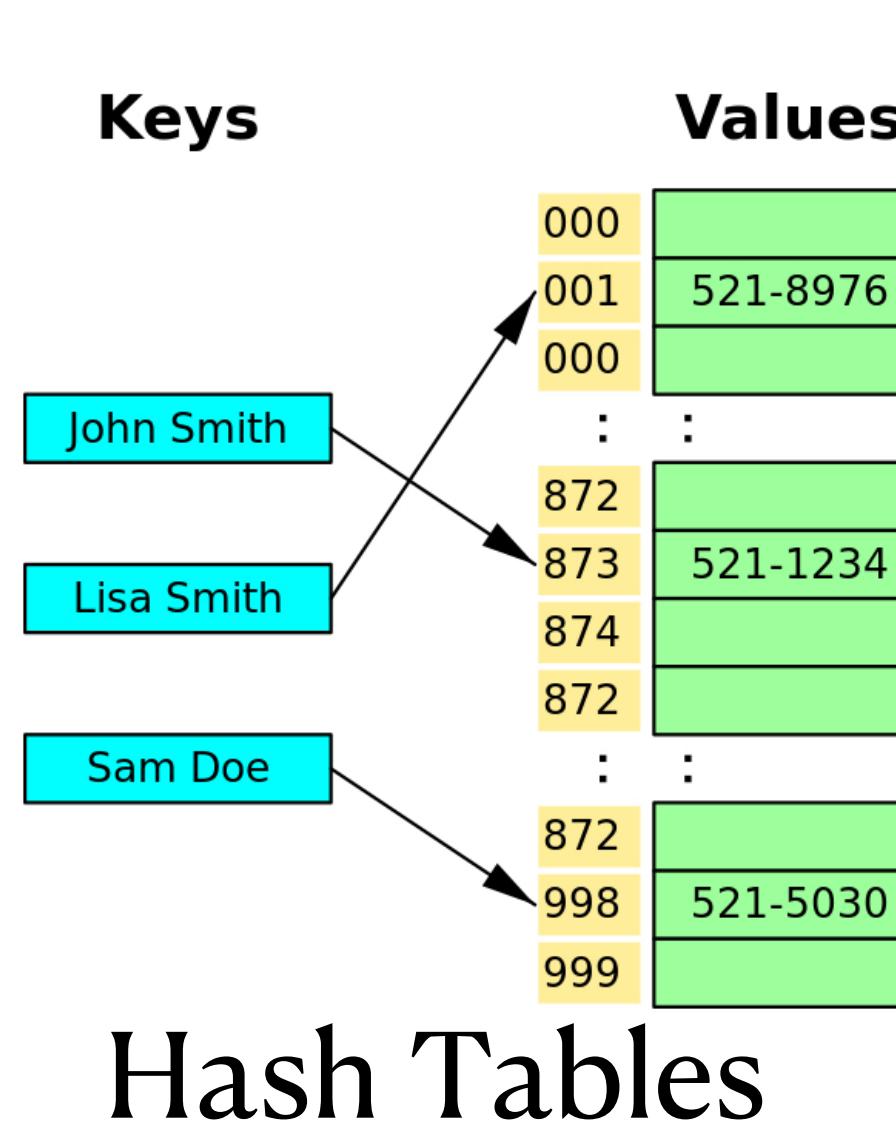
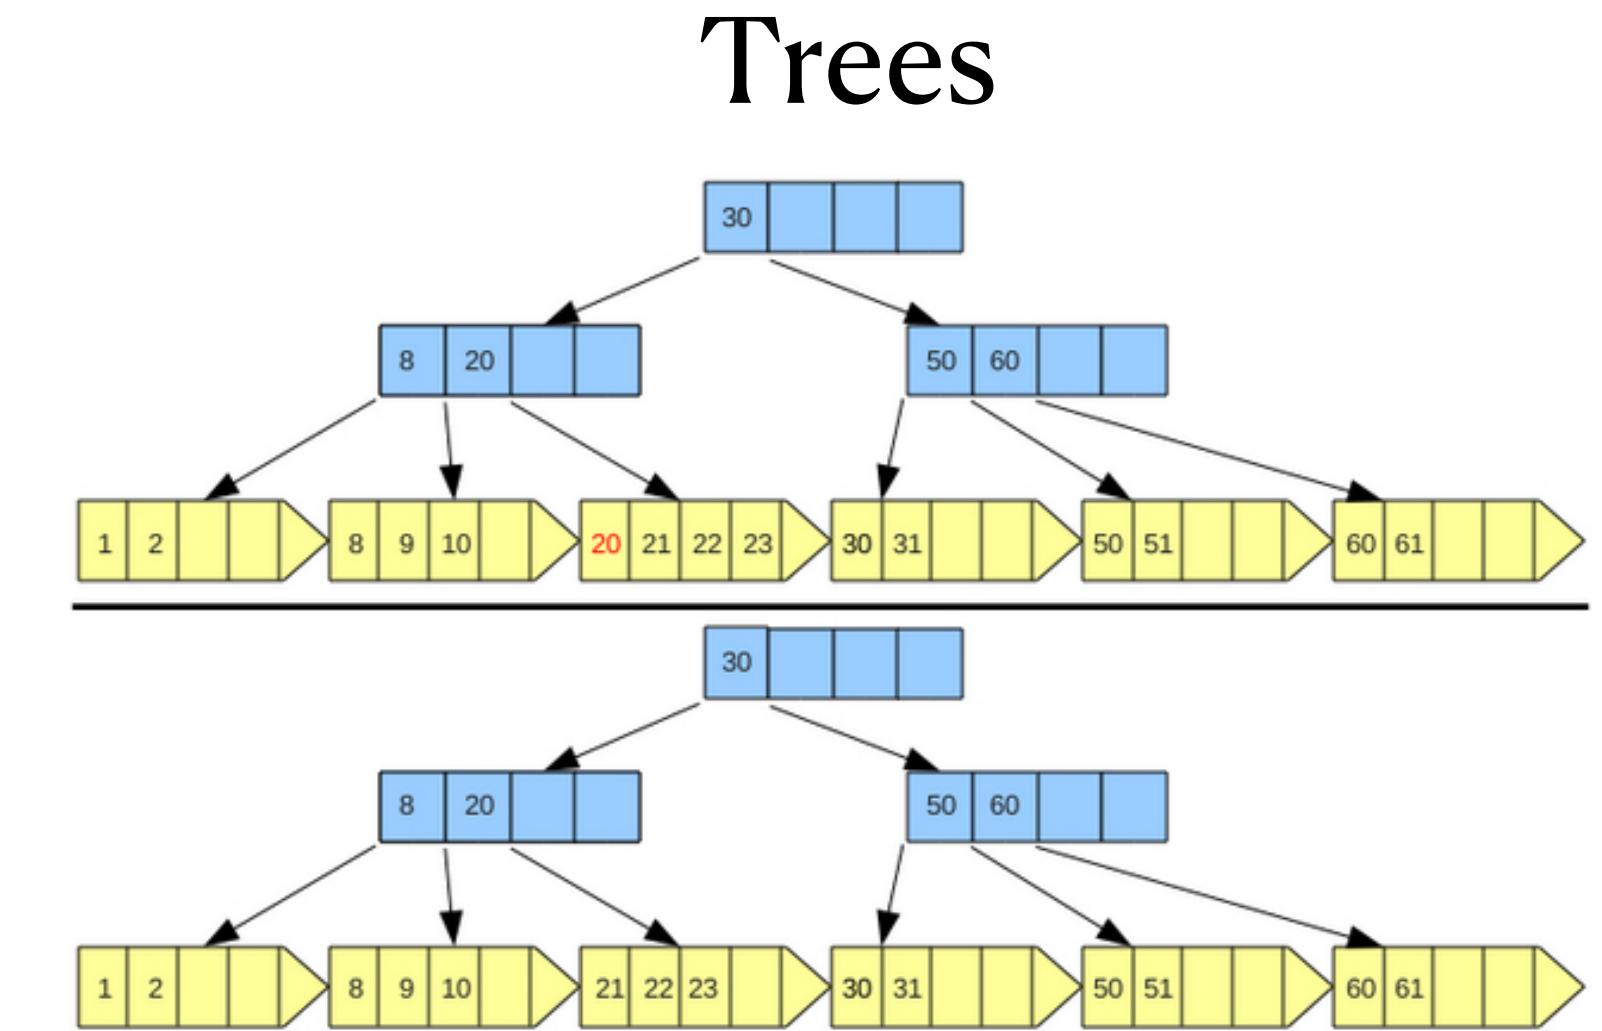
- Without indexing, searching data would require **full table scans**— painfully slow!
- Indexing **makes lookups fast** by keeping a structured map of data locations.



Indexing

Types of Indexes

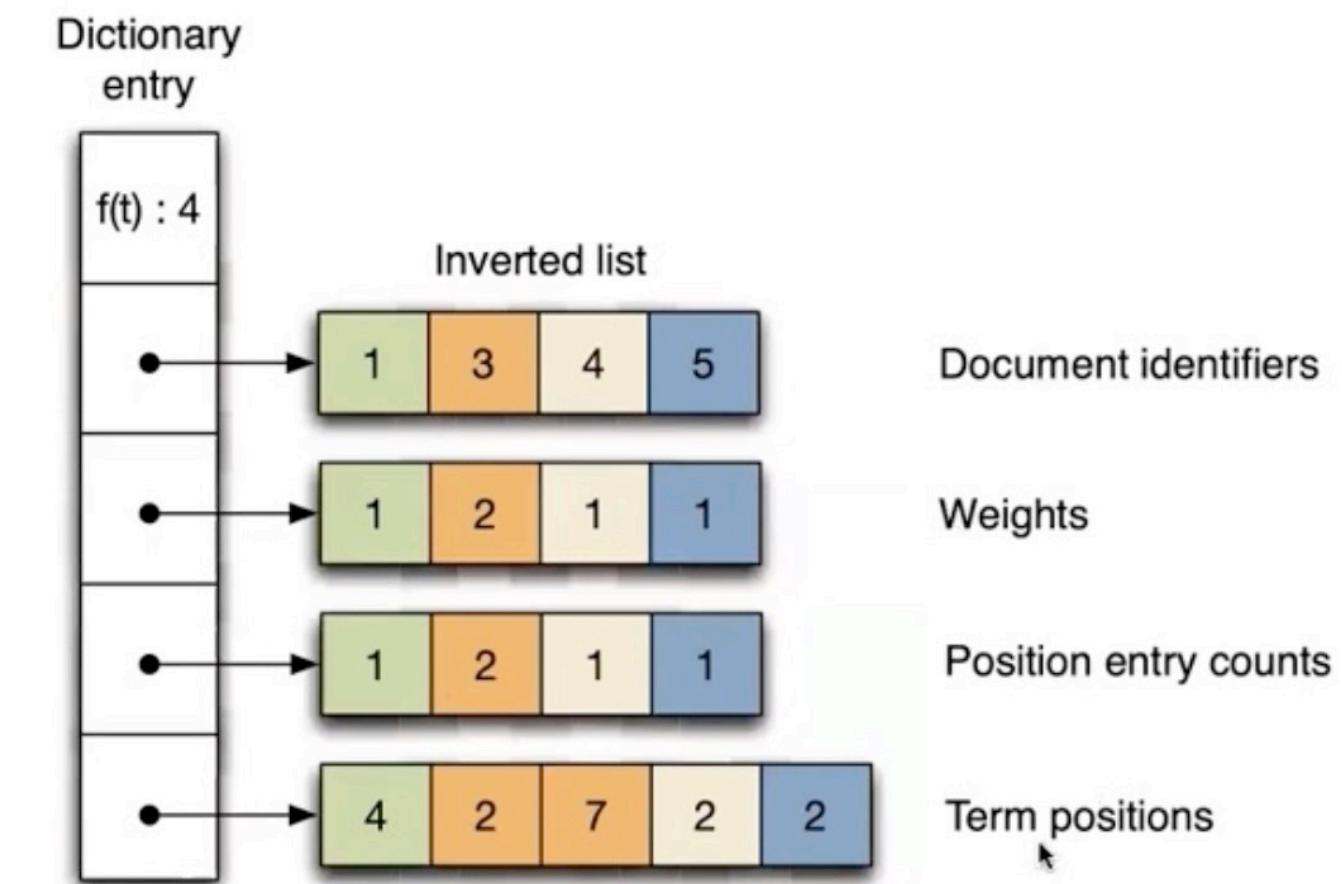
- B-Trees: Backbone of traditional indexing (used in almost all databases)
- Hash Indexes: Used for key-value stores
- LSM Trees: Used in write-heavy No-SQL systems
- Learned Indexes (?)



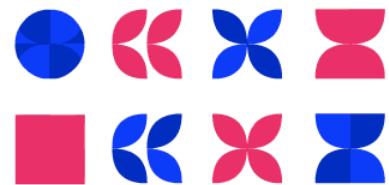
Indexing

Beyond Databases

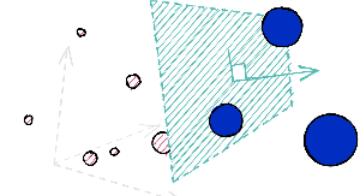
- Search engines – inverted index
- AI Retrieval – vector databases



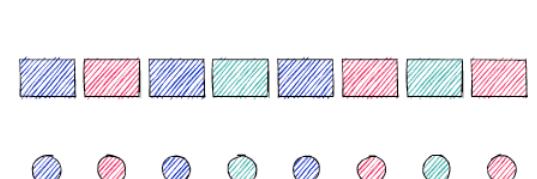
Similarity Search
TF-IDF | BM25 | Sentence-BERT



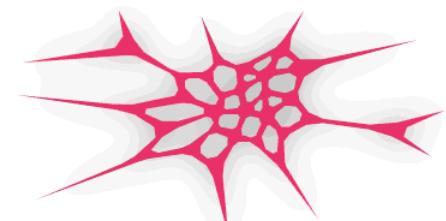
Similarity Search
LSH in Facebook AI Similarity Search



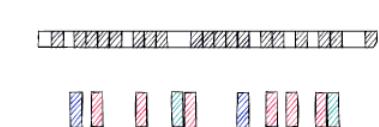
Similarity Search
Product Quantization 101



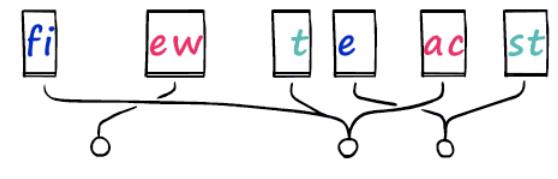
FAISS
Scalable Search With Facebook AI



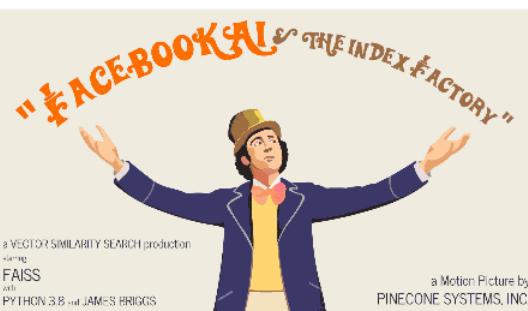
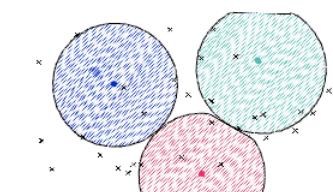
Similarity Search
The Missing WHERE Clause



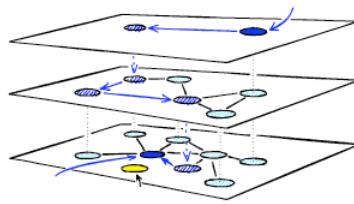
Similarity Search
Locality Sensitive Hashing



Similarity Search
Choosing the Right Index



Similarity Search
Hierarchical Navigable Small Worlds (HNSW)

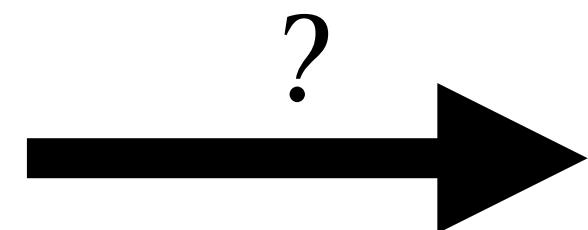


B. Cambazoglu and R. Baeza-Yates, "Scalability Challenges in Web Search Engines"

Query Optimization and Execution

Making declarative work imperatively

```
Select id, first_name, last_name  
from actor left join film on ...  
where exist (  
... )
```



?

```
HashTable<Integer, Record>  
hashTable = new HashTable...  
for(int i = 0; i < table1.numRows;  
i++) {  
    hashTable.put(table.getField(i));  
    ...  
} for(int i;i < table2.numRows; i++)  
{  
    ...  
    ...  
}
```

You describe what you want

The database figures out the
best way to execute your query

Query Optimization and Execution

Making declarative work imperatively

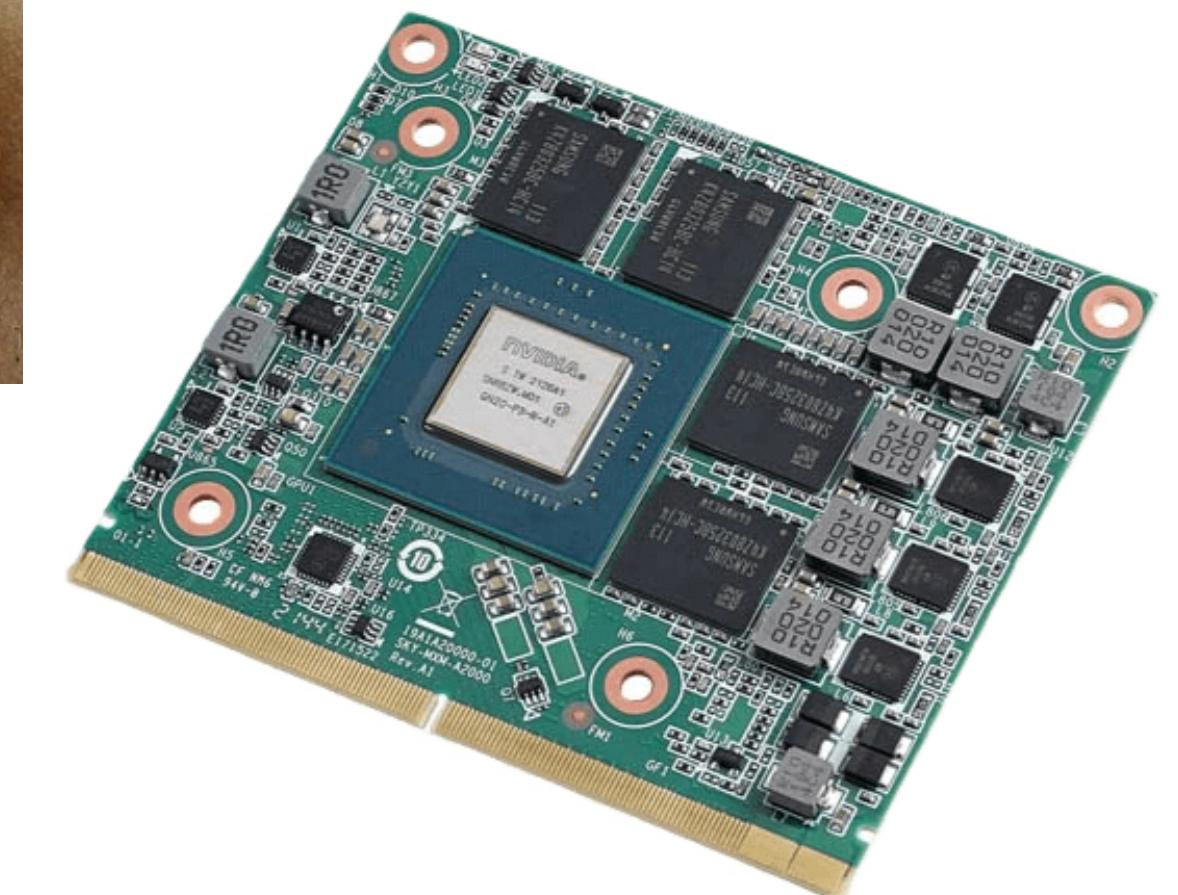
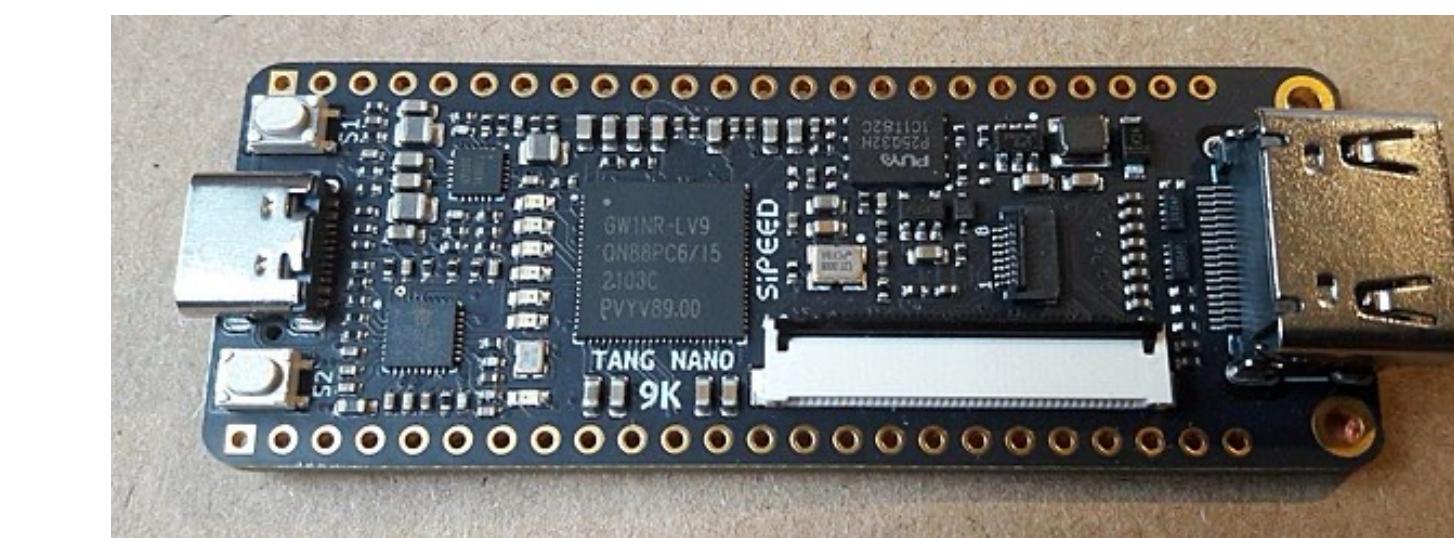
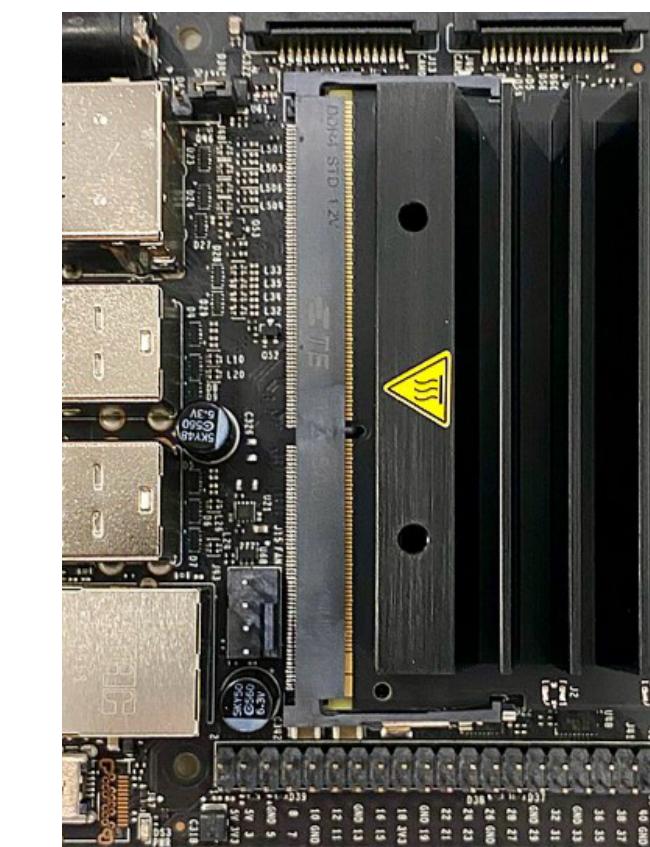
Key steps in query execution

1. Parsing and validation
2. Query optimization
 - a. Generate multiple possible execution plans
 - b. Use cost-based optimization to chose the best one
3. Execution: run the plan using techniques like parallel execution and caching

Query Optimization and Execution

Algorithms

1. Join
2. Scans
3. Sorting
4.



Aware of underlying compute and storage hardware

Why care about DB internals?

- Database internals are at the core of the biggest technological revolutions
- Databases are not just about SQL
- **AI & Machine Learning:** Data pipelines rely on indexing and query optimization.
- **Cloud Infrastructure:** Distributed databases like Spanner run on database internals.
- **Streaming & Real-time Analytics:** Kafka, Apache Flink use database-style event processing.
- **Graph Analytics:** Used in social networks, fraud detection, and knowledge graph

Every major tech company – Google, Amazon, Facebook – is, at its core, a database company